# ACCEPTED | Applicant 209765

SHAP and LIME explain the same LightGBM input and the same default-class output

**Model score 0.274**

Accept < 0.500 | Reject >= 0.500

## SHAP

Feature value + exact TreeSHAP attribution

Top-10 attribution coverage: 69.3% | Raw-input score: 0.454

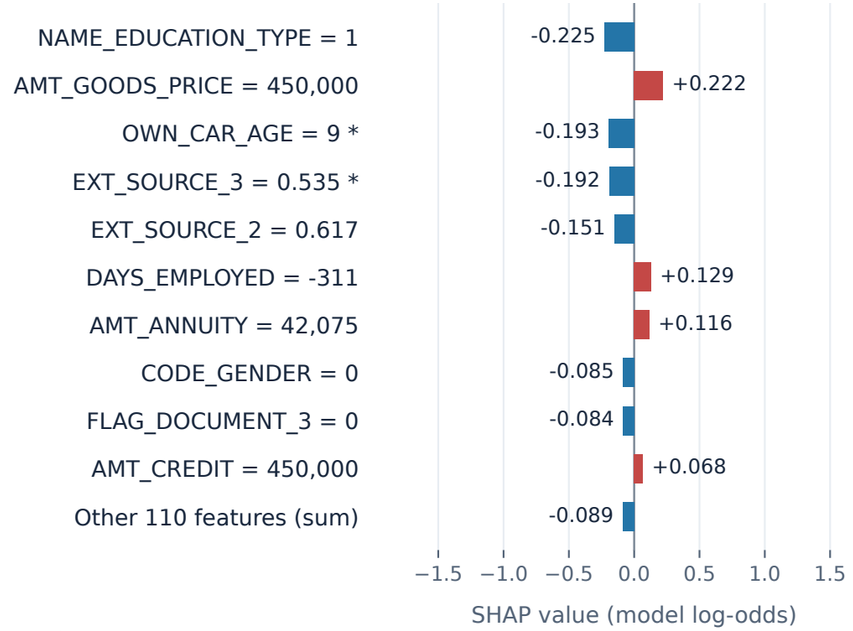


Blue / negative: lowers default-class score

## LIME

Discretized condition + local surrogate weight

Local fidelity $R^2$: 0.061 | 1,000 perturbations | seed 42

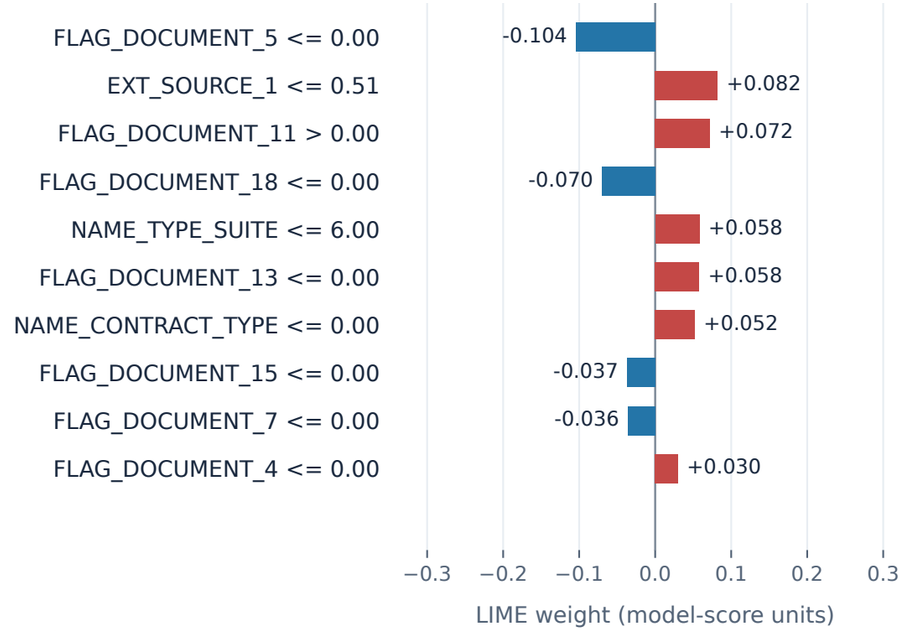


Red / positive: raises default-class score

**EXACT MODEL RECONSTRUCTION**

Base -0.493 + all contributions -0.483 = -0.976

Logistic transform -> 0.274 model score

**LOCAL APPROXIMATION CHECK**

Surrogate 0.397 | Model 0.274 | Gap 0.123

Condition boundaries describe bins; they are not approval thresholds.

Same median-filled input for both methods; * = imputed value. Top 10 ranked within each method. Panel units differ.

Recomputed from paper code. Acceptance/rejection is a model rule, not a recorded lending decision. Score is uncalibrated.